

Mynor A. Diaz

Broward County, Florida

<https://www.linkedin.com/in/mynor-diaz> | Portfolio: <https://madxxi.github.io/adam-portfolio/>

SYSTEMS AND SERVICE RELIABILITY LEADER

Senior Site Reliability Engineer | Cloud Infrastructure & Observability | Azure & Private Cloud

Senior SRE & Operations Engineer with 20+ years delivering high-availability platforms for Amadeus IT Group, one of the world's largest travel reservation and distribution networks, consistently achieving 98–99% availability across 460+ mission-critical applications by reducing operational toil through automation, observability tooling, and blameless incident management. Led alerts-as-code initiatives (over 100) at Amadeus using Prometheus, Thanos, and Grafana to standardize SLO-based monitoring at scale. Expertise also includes infrastructure-as-code practices developed through hands-on portfolio work in AWS and Terraform. Proven ability to own complex services end-to-end while driving efficiency, scalability, and reliability.

TECHNICAL SKILLS

Cloud & Container Platforms:	Azure, Kubernetes (v1.27+), Red Hat OpenShift (v4.x)
Cloud & IaC (Self-Directed/Portfolio Project):	AWS (EC2, S3, IAM), Terraform — hands-on via personal observability and infrastructure projects; <i>certifications in progress</i>
Observability & Monitoring:	Thanos, Grafana, Splunk, Prometheus (Datadog-equiv. metrics/alert stack)
Programming & Automation:	Go, Unix/Linux Shell (bash/ksh/sh), Perl, PHP, GitHub Actions, Jira, Python, Resolve
Incident, Problem & Change Management:	ITIL, Blameless Post-Incident Reviews, On-Call Rotation, Cap. Planning
Databases:	Supported: Oracle, MariaDB, MySQL. Usage of SQL, MySQL, promQL
Tools:	Git/Github, JIRA, VMware ESXi, Docker, Kubernetes, BIRD, Deployment Manager, Jenkins pipelines
Frameworks:	IT Services Management (ITSM), Scaled Agile Framework (SaFe)
Operating Systems:	Linux RHEL, SuSe Linux, CentOS, Ubuntu, AIX, Solaris, Mac OS, HP/UX
Documentation/Other:	Docker, Git, Jira/Confluence, BMC Control-M, Agile/SAFe, Sharepoint Yammer/wiki/jabber, Load Spikes Mgt, Large-Scale Distributed Systems

PROJECTS

- ***Spearheaded a JIRA EPIC to enhance, consolidate, and standardize Cloud alerts:***
 - Presented to the organization's Development Department a model for standardized alerts (templates) to be used by the entire support organization and helped maintain these to adapt them to all levels required.
 - Initiated coordination between development and operations teams to migrate alert events from internal "alerting system" into ServiceNow.
 - Pioneered, after "as-code" templates were created, to extend and amend standard alerts as needed (written in GO) using queries in PromQL via Thanos to trigger those alerts and building on-the-fly Grafana dashboards.
- ***Elemental on the support of customer application and data migrations in Amadeus and all prior roles:***
 - 2018-2026: From physical and virtual Linux and windows systems to Cloud platforms 2018. Migrations to Openshift server 3.x and 4.x on private platforms followed by migration to public cloud: Azure.
 - 2005-2016: Migrations from Mainframes (IBMs TPF) to Unix/Linux Physical and VMWare + Veritas Cluster Services (VCS).

PROFESSIONAL EXPERIENCE

AMADEUS Global Operations Americas Inc. (Remote: 3/20-5/23 / Hybrid: 6/23-3/26)

Senior Service Reliability Engineer (SRE) - Cloud Technologies

Jan 2019 – Mar 2026

- Owned reliability and operations for 460+ mission-critical production applications (reservations, ticketing, inventory) serving global travel customers, consistently achieving 98–99% availability and strict SLOs in a 24/7 high-traffic SaaS environment. Applied SLI/SLO frameworks and error-budget-based burn-rate alerting to prioritize reliability work and reduce false-positive incident noise across thousands of monitored metrics.
- Spearheaded a major JIRA Epic to standardize, consolidate, and automate alerts using Prometheus, Thanos, PromQL, Go, and ServiceNow. ***With over 100 template-based alerts developed***, significantly reduced alert noise, enabled on-the-fly Grafana dashboards, and improved operational efficiency across the organization.

- Led large-scale cloud migrations from legacy physical/virtual Linux systems to private OpenShift (v3/v4) and public Azure, delivering zero business disruption while modernizing infrastructure and reducing operational overhead.
- **Drove end-to-end incident, problem, and change management using ITIL best practices:** championed blameless postmortems and continuous improvements that enhanced system stability and recovery speed.
- **Developed automation scripts and tooling** (Go, Shell, GitHub Actions) to reduce manual toil, accelerate deployments, and improve monitoring and recovery processes.
- Have applied AIOps practices (AI-assisted troubleshooting with Copilot, Claude, GPT-4) and RAG-based retrieval of internal documentation (Confluence, technical procedures, OPROCs, etc.) to reduce operational toil and average troubleshooting time on complex incidents and accelerating runbook generation for new alert types (ServiceNow).
- **Led cloud technology adoption and large-scale migrations from legacy physical/virtual Linux systems to private OpenShift (v3.x/v4.x) and public Azure platforms** (2020–2021), while maintaining zero business disruption.
- **Trained and onboarded 25+ engineers** as the official instructor for the Amadeus Cloud Services (ACS) curriculum, delivering core courses on “ACS Induction” and “ACS Incidents and Recovery” within an Agile SAFe 5.x framework.
- **Built and fine-tuned advanced SRE tooling using Prometheus, Thanos, Grafana, PromQL, and GitHub Actions;** created custom operational procedures and integrated alerting with Microsoft Teams and ServiceNow during the full migration to a Prometheus-based architecture.

Amadeus North America

Lead Systems Engineer (UNIX/Linux),

2008 - 2018

Sr. Systems Engineer,

2004 - 2008

- **Served as technical lead for a global 24/7 team of 9 SREs** coaching, mentoring, and directing day-to-day technical work without formal reporting authority, while providing significant input into hiring and personnel decisions and driving mission-critical platform availability to an average of 99% for systems supporting the global travel industry.
- Directed all incident, change, and problem management using ITIL best practices.
- **Provided technical performance assessments and input into formal review cycles for engineers on the team,** directly informing professional development and advancement decisions, without holding formal reporting authority.
- **Served as Incident Commander / Incident Coordinator** (2013–2018) for major production incidents, directing cross-team response, driving real-time decision-making, and owning incident communications through to resolution.
- **Directed end-to-end major incident, problem, and change management** following ITIL best practices, driving root cause identification and preventive measures across thousands of production systems.
- **Owned operational stability and risk management for a multi-million Euro environment** spanning over 5,500 production and test systems (HP-UX, Linux, multiple architectures), overseeing high-availability cluster operations (Veritas, MC/ServiceGuard) and all major OS/database upgrades.
- **Championed automation and reliability engineering,** developing custom scripts and tools (Perl, PHP, HTML) to enhance monitoring, accelerate fault recovery, and eliminate toil, directly supporting the team’s ability to meet stringent availability targets.
- **Managed cross-functional technical execution,** including Oracle version migrations, IBM MQ incident resolution, and oversight of administrative environments (Dev/Pre-Prod/Prod), ensuring coordinated, low-risk change implementation across the entire technology stack.

EARLY CAREER HIGHLIGHTS

- UNIX System Administrator, Rollins, Inc., Atlanta, GA Led high-impact production migration of national call-center environment from legacy HP-UX/Oracle to new platform without external contractors, delivering immediate stability improvements and \$10,000 in cost savings.
- Senior Personnel System & Database Administrator, U.S. Army, Fort Benning, GA Directed 5-person team supporting 19 servers; designed and streamlined network infrastructure that reduced operational effort by 14,000 man-hours per year and generated \$200,000 in resource-management savings. Enhanced division-level database prototype by 40%, driving \$28 million in annual Army-wide cost savings while training 200+ end users.

EDUCATION & CERTIFICATIONS

- **Masters of Executive Management in Business Administration (EMBA):** University of Florida, 2013
- **Masters of Computer Information Systems (MIS):** Nova Southeastern University, FL, 2007
- **Certifications:** AZ Fundamentals, FinOps Foundation, MS AZ Data Fundamentals, SAFe 5 Practitioner, ITIL. Pursuing: HashiCorp Terraform Associate, AWS Certified SysOps Administrator