

Mynor A. Diaz

Weston, Florida

305-766-8750 | <https://www.linkedin.com/in/mynor-diaz> | Portfolio: <https://madxxi.github.io/adam-portfolio/>

SYSTEMS AND SERVICE RELIABILITY LEADER

Senior Site Reliability Engineer | Cloud Infrastructure & Observability | Azure & Private Cloud

Senior Site Reliability Engineer with 20+ years at Amadeus IT Group, one of the world's largest travel reservation networks, owning reliability for 460+ mission-critical applications. Sustained a 196–200/200 Service Stability Index for 12 consecutive years with sub-20-minute MTTR on critical outages, and served as Incident Commander for ~36 Sev-1/Sev-2 incidents per year. Led an alerts-as-code initiative (100+ Prometheus/Thanos rules, Go tooling, Grafana) that cut alert noise 50%, and the zero-downtime migration of 400+ applications to Red Hat OpenShift and Microsoft Azure. Extending into AWS and Terraform through hands-on portfolio work.

TECHNICAL SKILLS

Cloud & Container Platforms:	Azure, Kubernetes (v1.27+), Red Hat OpenShift (v4.x)
Cloud & IaC (Self-Directed/Portfolio Project):	AWS (EC2, S3, IAM), Terraform — hands-on via personal observability and infrastructure projects; <i>certifications in progress</i>
Observability & Monitoring:	Thanos, Grafana, Splunk, Prometheus (Datadog-equiv. metrics/alert stack)
Programming & Automation:	Go, Unix/Linux Shell (bash/ksh/sh), Perl, PHP, GitHub Actions, Jira, Python, Resolve
Incident, Problem & Change Management:	ITIL, Blameless Post-Incident Reviews, On-Call Rotation, Cap. Planning
Databases:	Supported: Oracle, MariaDB, MySQL. Usage of SQL, MySQL, promQL
Tools:	Git/Github, JIRA, VMware ESXi, Docker, Kubernetes, BIRD, Deployment Manager, Jenkins pipelines
Frameworks:	IT Services Management (ITSM), Scaled Agile Framework (SaFe)
Operating Systems:	Linux RHEL, SuSe Linux, CentOS, Ubuntu, AIX, Solaris, Mac OS, HP/UX
Documentation/Other:	Docker, Git, Jira/Confluence, BMC Control-M, Agile/SAFe, Sharepoint Yammer/wiki/jabber, Load Spikes Mgt, Large-Scale Distributed Systems

PROJECTS

- ***Spearheaded a JIRA EPIC to enhance, consolidate, and standardize Cloud alerts:***
 - Presented to the organization's Development Department a model for standardized alerts (templates) to be used by the entire support organization and helped maintain these to adapt them to all levels required.
 - Initiated coordination between development and operations teams to migrate alert events from internal "alerting system" into ServiceNow.
 - Pioneered, after "as-code" templates were created, to extend and amend standard alerts as needed (written in GO) using queries in PromQL via Thanos to trigger those alerts and building on-the-fly Grafana dashboards.
- ***Elemental on the support of customer application and data migrations in Amadeus and all prior roles:***
 - 2018-2026: From physical and virtual Linux and windows systems to Cloud platforms 2018. Migrations to Openshift server 3.x and 4.x on private platforms followed by migration to public cloud: Azure.
 - 2005-2016: Migrations from Mainframes (IBMs TPF) to Unix/Linux Physical and VMWare + Veritas Cluster Services (VCS).

PROFESSIONAL EXPERIENCE

AMADEUS Global Operations Americas Inc.

Remote: 3/20-5/23 / Hybrid: 6/23-3/26

Senior Service Reliability Engineer (SRE) - Cloud Technologies

Jan 2019 – Mar 2026

- Maintained a 196–200/200 Service Stability Index (SSI), Amadeus's internal reliability scorecard tied to MTTR, for 12 consecutive years (2013–2025), resolving 90%+ of incidents within target: sub-20-minute MTTR on critical outages and sub-30-minute on degradations.
- Implemented SLI/SLO frameworks and multi-window error-budget burn-rate alerting across 460+ mission-critical travel applications, cutting false-positive pages by 50% and focusing engineering effort on the highest-risk reliability work.

- Built a library of 100+ templated Prometheus/Thanos (PromQL) alert rules with Go-based tooling that auto-generated Grafana dashboards and routed incidents to ServiceNow; automated rule rollout with GitHub Actions and Bash, reducing on-call pages by ~250/week, deployment time from 120 to 15 minutes, and manual toil by ~10 hours/week.
- Served as Incident Commander for ~36 Sev-1/Sev-2 incidents per year, coordinating 10–20+ interdependent application teams (inventory, reservations, ticketing, pricing, baggage) plus network, storage, infrastructure, and Microsoft vendor teams per incident under ITIL incident, problem, and change management; ran a blameless postmortem for every Sev-1/Sev-2, with corrective actions that reduced repeat incidents by 20% and lowered change failure rate from 15% to 1%.
- Reduced average troubleshooting time on complex incidents by 15% by introducing AIOps practices, including AI-assisted diagnosis (GitHub Copilot, Claude, GPT-4) and RAG-based retrieval over Confluence runbooks and operational procedures; cut runbook creation for new ServiceNow alert types from 2 hours to 15 minutes.
- Led the multi-phase migration of 400+ applications and 2,600 Linux hosts from physical to virtual infrastructure, then to private Red Hat OpenShift (v3/v4) and public Microsoft Azure, decomposing the portfolio into 460+ microservices-based applications with zero customer-facing downtime and reducing provisioning time from 6 weeks to 48 hours.
- Trained and onboarded 25+ engineers as official instructor for the Amadeus Cloud Services (ACS) curriculum (ACS Induction; ACS Incidents and Recovery), shortening new-hire time to operational readiness from 12 to 2 weeks.

Amadeus North America

Lead Systems Engineer (UNIX/Linux),

2008 - 2018

Sr. Systems Engineer,

2005 - 2008

- Served as technical lead for a global 24/7 team of 9 engineers without formal reporting authority, providing performance assessments that fed formal review cycles, and significantly influencing hiring decisions; mentored 18 engineers into senior roles over 10 years and cut team turnover by 25%.
- Served as Incident Commander (2013–2018) for ~15 major production incidents per year, directing cross-team response and stakeholder communications through resolution; led ITIL incident, problem, and change management with root-cause and preventive actions across thousands of production systems, reducing repeat incidents by 95% and change-related outages by 90%.
- Owned operational stability and risk management for a €2.7B revenue-supporting environment of 5,500+ production and test systems (HP-UX, Linux), overseeing high-availability clusters (Veritas, MC/ServiceGuard) and leading 250 major OS and database upgrades while sustaining 99.99% uptime on clustered production services.
- Developed Perl and PHP monitoring and recovery tools that eliminated ~20 hours/week of manual checks; the tooling was showcased at sites in France and England and led to the formation of a dedicated alerting and monitoring team in 2014.
- Supported 10 Oracle version migrations and IBM MQ incident resolution across Dev, Pre-Prod, and Prod environments, completing all within planned maintenance windows with zero data loss.

EARLY CAREER HIGHLIGHTS

- UNIX System Administrator, Rollins, Inc., Atlanta, GA Led high-impact production migration of national call-center environment from legacy HP-UX/Oracle to new platform without external contractors, delivering immediate stability improvements and \$10,000 in cost savings.
- Senior Personnel System & Database Administrator, U.S. Army, Fort Benning, GA Directed 5-person team supporting 19 servers; designed and streamlined network infrastructure that reduced operational effort by 14,000 man-hours per year and generated \$200,000 in resource-management savings. Enhanced division-level database prototype by 40%, driving \$28 million in annual Army-wide cost savings while training 200+ end users.

EDUCATION & CERTIFICATIONS

- **Masters of Executive Management in Business Administration (EMBA):** University of Florida, 2013
- **Masters of Computer Information Systems (MIS):** Nova Southeastern University, FL, 2007
- **Certifications:** AZ Fundamentals, FinOPs Foundation, MS AZ Data Fundamentals, SAFe 5 Practitioner, ITIL. Pursuing: HashiCorp Terraform Associate, AWS Certified SysOps Administrator