

Mynor A. Diaz

Florida

mr.mynor.diaz@gmail.com | <https://www.linkedin.com/in/mynor-diaz>

Portfolio: <https://madxxi.github.io/adam-portfolio/>

Lead SRE | Cloud Infrastructure & Observability | Major Incident Response

Senior SRE & Operations Engineer with 20+ years supporting high-availability platforms for Amadeus IT Group, one of the world's largest travel reservation and distribution networks, consistently achieving 98–99% availability across 460+ mission-critical applications by reducing operational toil through automation, observability tooling, and blameless incident management. Led alerts-as-code initiatives at Amadeus using Prometheus, Thanos, and Grafana to standardize SLO-based monitoring at scale. Expertise with infrastructure-as-code practices developed through hands-on portfolio work in AWS and Terraform. Proven ability to own complex/distributed global production highly available critical services end-to-end while driving efficiency, scalability, and reliability. Maintained a 196-200 out of 200 Service Stability Index (SSI), Amadeus's internal reliability scorecard tied directly to MTTR performance, for 12 consecutive years (2013–2025), reflecting sustained sub-20-minute MTTR on critical outages and sub-30-minute MTTR on degradations across 90%+ of incidents.

PROJECTS

- ***Originated and led an organization-wide reliability initiative to standardize cloud alerts:***
 - Presented to the organization's Development Department a model for standardized alerts (templates) to be used by the entire support organization and helped maintain these to adapt them to all levels required.
 - Spearheaded coordination to migrate alert events from internal "alerting system" into ServiceNow.
 - Pioneered "as-code" templates to create standard alerts as needed (written in GO) with PromQL via Thanos.

EXPERIENCE

AMADEUS Global Operations Americas Inc.

Sr. Service Reliability Engineer (SRE) - Cloud Technologies

2019 – 2026

Fully Remote Mar 2020 – May 2023 / Hybrid Jun 2023 - Mar 2026

- **Led enterprise-wide JIRA Epic that standardized template-based monitoring alerts** (written in Go with PromQL/Thanos), integral part on migration of events into ServiceNow, reduced alert noise/duplication by eliminating redundant alerts, and delivered on-the-fly Grafana dashboards, now adopted across the entire organization for 460+ global travel applications.
- Directed end-to-end reliability engineering for 460+ mission-critical production applications (reservations, ticketing, inventory, loyalty) in a 24/7 high-availability SaaS environment, consistently achieving and sustaining 98–99% uptime while meeting strict SLAs and SLOs for global B2B customers (Airlines, Hotels, Cruise Lines, Car Rental, etc.).
- **Spearheaded proactive observability and capacity planning initiatives**, implementing autoscaling, advanced monitoring strategies, and alerting solutions that prevented outages during extreme seasonal load spikes and major travel events.
- **Served as senior leader on the Incident & Problem Management team**, driving rapid recovery of critical services, minimizing customer impact, and championing blameless postmortems and continuous process improvements to the Change Management framework over a decade.
- **Pioneered changes for transition from proprietary/internal event-management tooling** to industry standard ServiceNow for all platform alerts, significantly improving operational efficiency and reducing false positives.
- **Built and fine-tuned advanced SRE tooling using Prometheus, Thanos, Grafana, PromQL, and GitHub Actions**; created custom operational procedures and integrated alerting with Microsoft Teams and ServiceNow during the full migration to a Prometheus-based architecture.
- Applied SRE and observability principles across Azure/OpenShift environments.
- Used AI coding assistants (GitHub Copilot, Claude, GPT-4) and RAG-based retrieval against internal documentation (Confluence, OPROCs) to accelerate troubleshooting and remediation via internal ticketing system, SNOW, Thanos alerts.
- **Led cloud technology adoption and large-scale migrations from legacy physical/virtual Linux systems to private OpenShift (v3.x/v4.x) and public Azure platforms** (2020–2021), while maintaining zero business disruption.
- **Trained and onboarded 25+ engineers** as the official instructor for the Amadeus Cloud Services (ACS) curriculum, delivering core courses on "ACS Induction" and "ACS Incidents and Recovery" within an Agile SAFe 5.x framework.

Amadeus North America,**Sr. Systems Engineer & Lead Systems Engineer (UNIX/Linux)****2004 – 2008 / 2008 - 2019**

- **Served as technical lead for a global 24/7 team of 9 SREs** coaching, mentoring, and directing day-to-day technical work without formal reporting authority, while providing significant input into hiring and personnel decisions and driving mission-critical platform availability to an average of 99% for systems supporting the global travel industry. Directed all incident, change, and problem management using ITIL best practices.
- **Served as Incident Commander / Incident Coordinator** (2013–2018) for major production incidents, directing cross-team response, driving real-time decision-making, and owning incident communications through to resolution.
- **Directed end-to-end major incident, problem, and change management** following ITIL best practices, driving root cause identification and preventive measures across thousands of production systems.
- **Advised senior management** on tooling, equipment, and team capacity investments, directly informing decisions on technology procurement and team expansion for growth and succession planning.
- **Played a central role in the technical interview process** for SRE and engineering hires, with assessments carrying substantial weight in final hiring decisions.
- **Owned operational stability and risk management for a multi-million Euro environment** spanning over 5,500 production and test systems (HP-UX, Linux, multiple architectures), overseeing high-availability cluster operations (Veritas, MC/ServiceGuard) and all major OS/database upgrades.
- **Championed automation and reliability engineering**, developing custom scripts and tools (Perl, PHP, HTML) to enhance monitoring, accelerate fault recovery, and eliminate toil, directly supporting the team's ability to meet stringent availability targets.
- **Managed cross-functional technical execution**, including Oracle version migrations, IBM MQ incident resolution, and oversight of administrative environments (Dev/Pre-Prod/Prod), ensuring coordinated, low-risk change implementation across the entire technology stack.

EARLY CAREER HIGHLIGHTS

- UNIX System Administrator, Rollins, Inc., Atlanta, GA Led high-impact production migration of national call-center environment from legacy HP-UX/Oracle to new platform without external contractors, delivering immediate stability improvements and \$10,000 in cost savings.
- Senior Personnel System & Database Administrator, U.S. Army, Fort Benning, GA Directed 5-person team supporting 19 servers; designed and streamlined network infrastructure that reduced operational effort by 14,000 man-hours per year and generated \$200,000 in resource-management savings. Enhanced division-level database prototype by 40%, driving \$28 million in annual Army-wide cost savings while training 200+ end users.

EDUCATION**University of Florida (2013): Masters of Executive Management in Business Administration (EMBA)****Nova Southeastern University, FL (2007): Masters of Computer Information Systems (MIS)****LICENSES & CERTIFICATIONS**

- *MS Azure Fundamentals*
- *FinOps Foundation*
- *Pursuing: HashiCorp Terraform Associate, AWS Certified SysOps Administrator*
- *MS Azure Data Fundamentals*
- *SAFe 5 Practitioner*

SKILLS

- **Cloud Platforms:** Azure & RH Openshift : Kubernetes (v1.27.x)/ OpenShift (v4.1x)
- **Cloud & IaC (Self-Directed / Portfolio Projects):** AWS (EC2, S3, IAM), Terraform, hands-on via personal observability and infrastructure projects; certifications in progress
- **Incident, Problem & Change Management:** ITIL, Blameless Post-Incident Reviews, On-Call Rotation, Cap. Planning
- **Databases:** Supported: Oracle, MariaDB,MySQL. Usage of SQL, MySQL, promQL
- **Tools:** Git/Github, JIRA, VMware ESXi, Docker, Kubernetes, BIRD, Deployment Manager, Jenkins pipelines
- **Frameworks:** IT Services Management (ITSM), Scaled Agile Framework (SaFe)
- **Observability and Monitoring:** Thanos, Grafana, Splunk, Prometheus (Datadog-equivalent metrics/alerting stack)
- **Automation & Alerting:** ServiceNow (Service Mgt, incident escalation), Resolve (automation), Jira (Agile)
- **Operating Systems:** Linux RHEL, SuSe Linux, CentOS, Ubuntu, AIX, Solaris, Mac OS, HP/UX Windows
- **Languages:** Unix / Linux Shells Scripts, sed, Awk, PHP, Python, Perl, GO
- **Communications and Documentation:** Atlassian: Confluence & JIRA / MS: SharePoint, Yammer / wiki, Jabber